Dataset distribution before modelling

(a) Class balance per split (IMDB, grouped by cover)

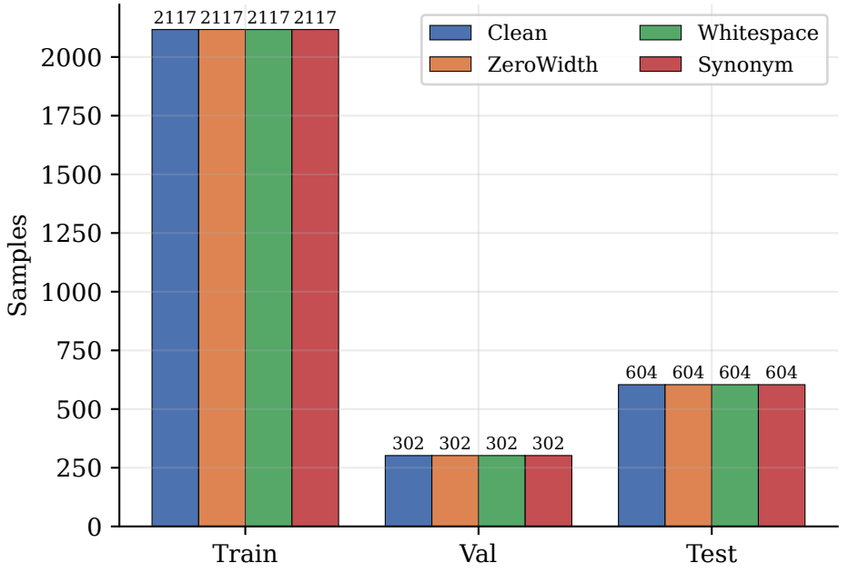

(b) Cover-text length by corpus

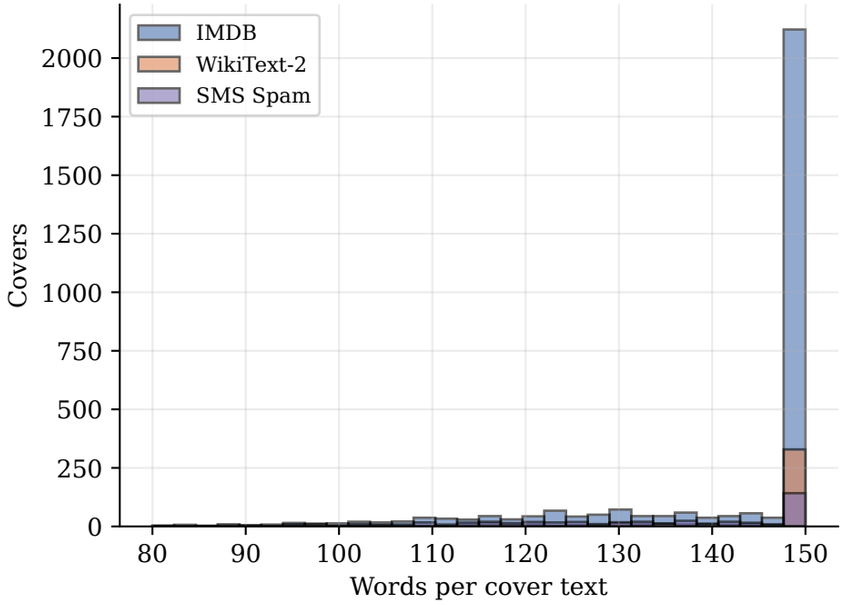

(c) Zero-width chars per sample

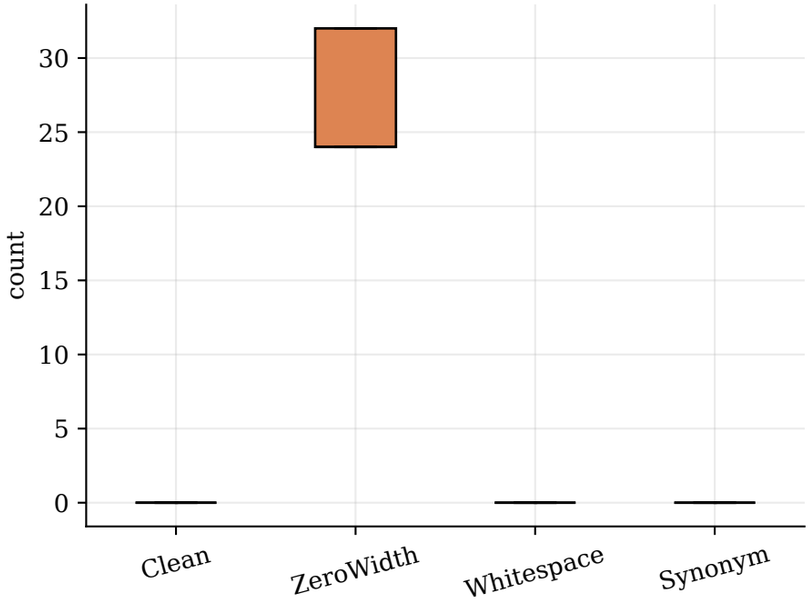

(d) Double spaces per sample

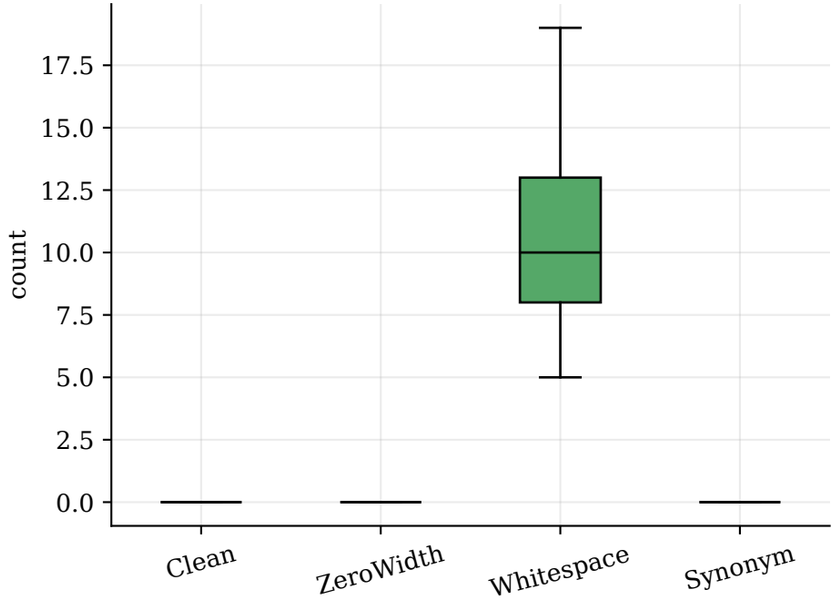

(e) Rare-synonym ratio (WordNet count = 0)

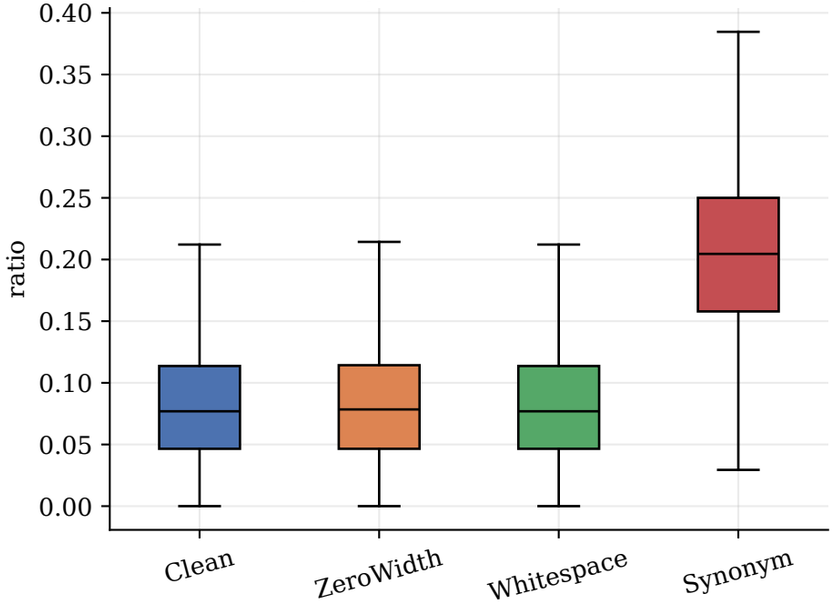

(f) Corpora summary (4 classes per cover)

Corpus	Role	Covers	Samples	Avg words
IMDB	train / val / test	3,023	12,092	142
WikiText-2	external val #1	476	1,904	142
SMS Spam	external val #2	475	1,900	130