Modern baselines (larger encoder and LoRA-tuned LLM) under the same protocol

(a) Macro-F1 (single seed; see the multi-seed figure for seed spread)

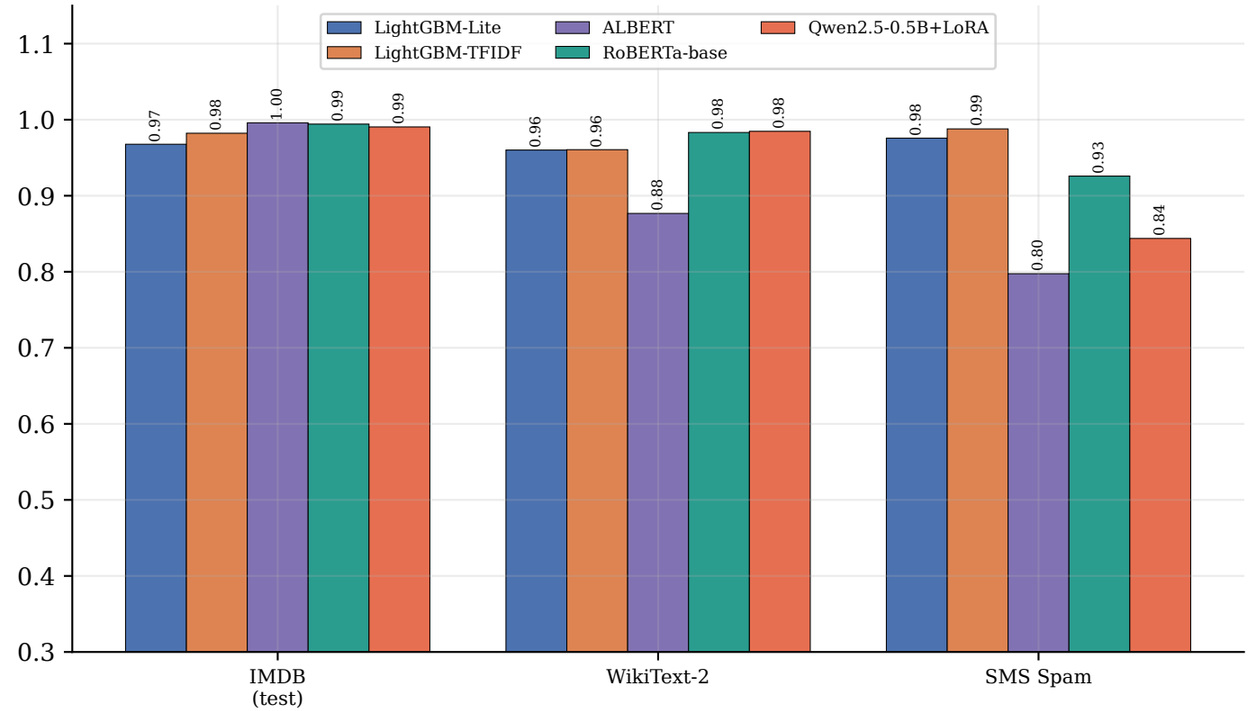

(b) TPR at the IMDB-val 1%-FPR threshold

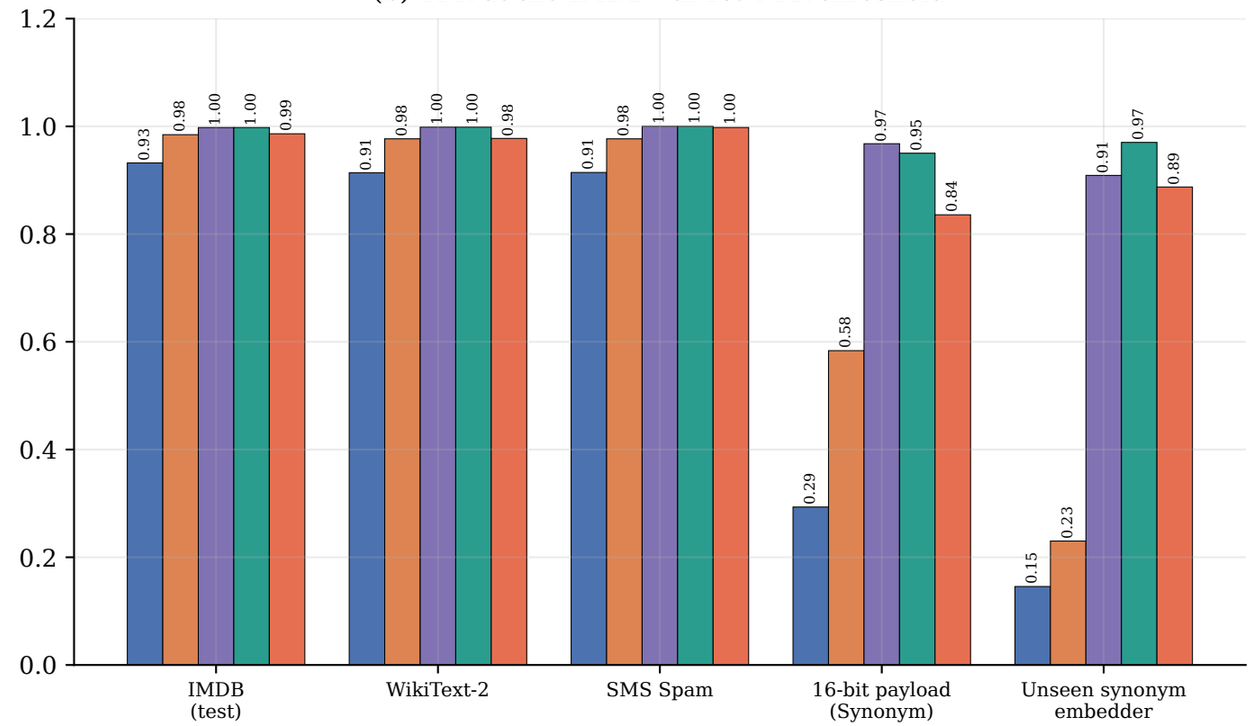

(c) FPR at the same threshold (target 1%)

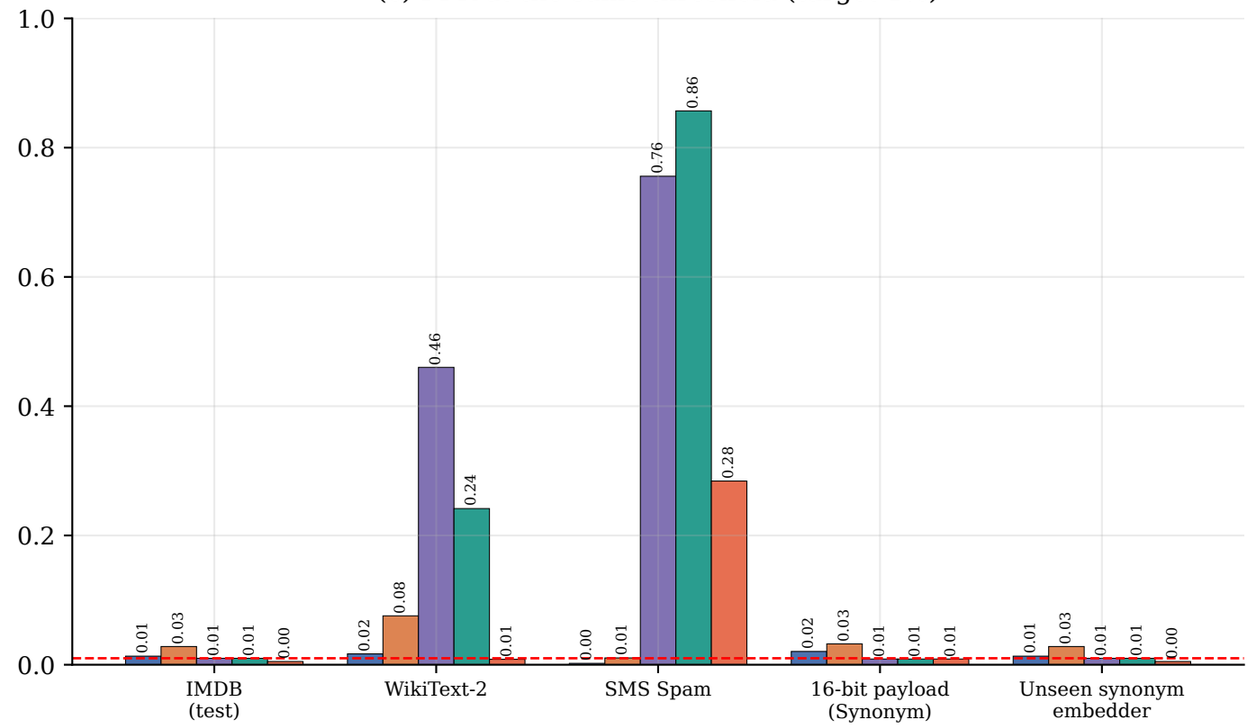
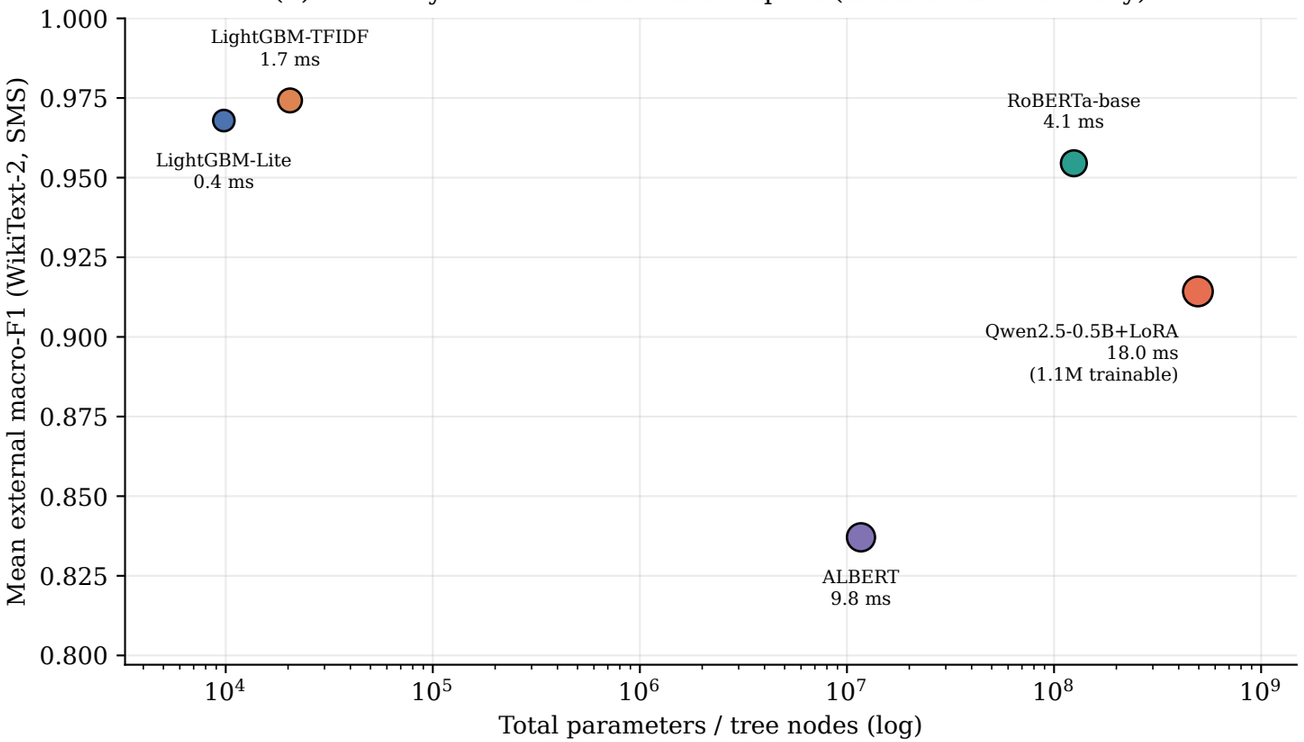
(d) Accuracy vs. cost on external corpora (marker size ~ latency)