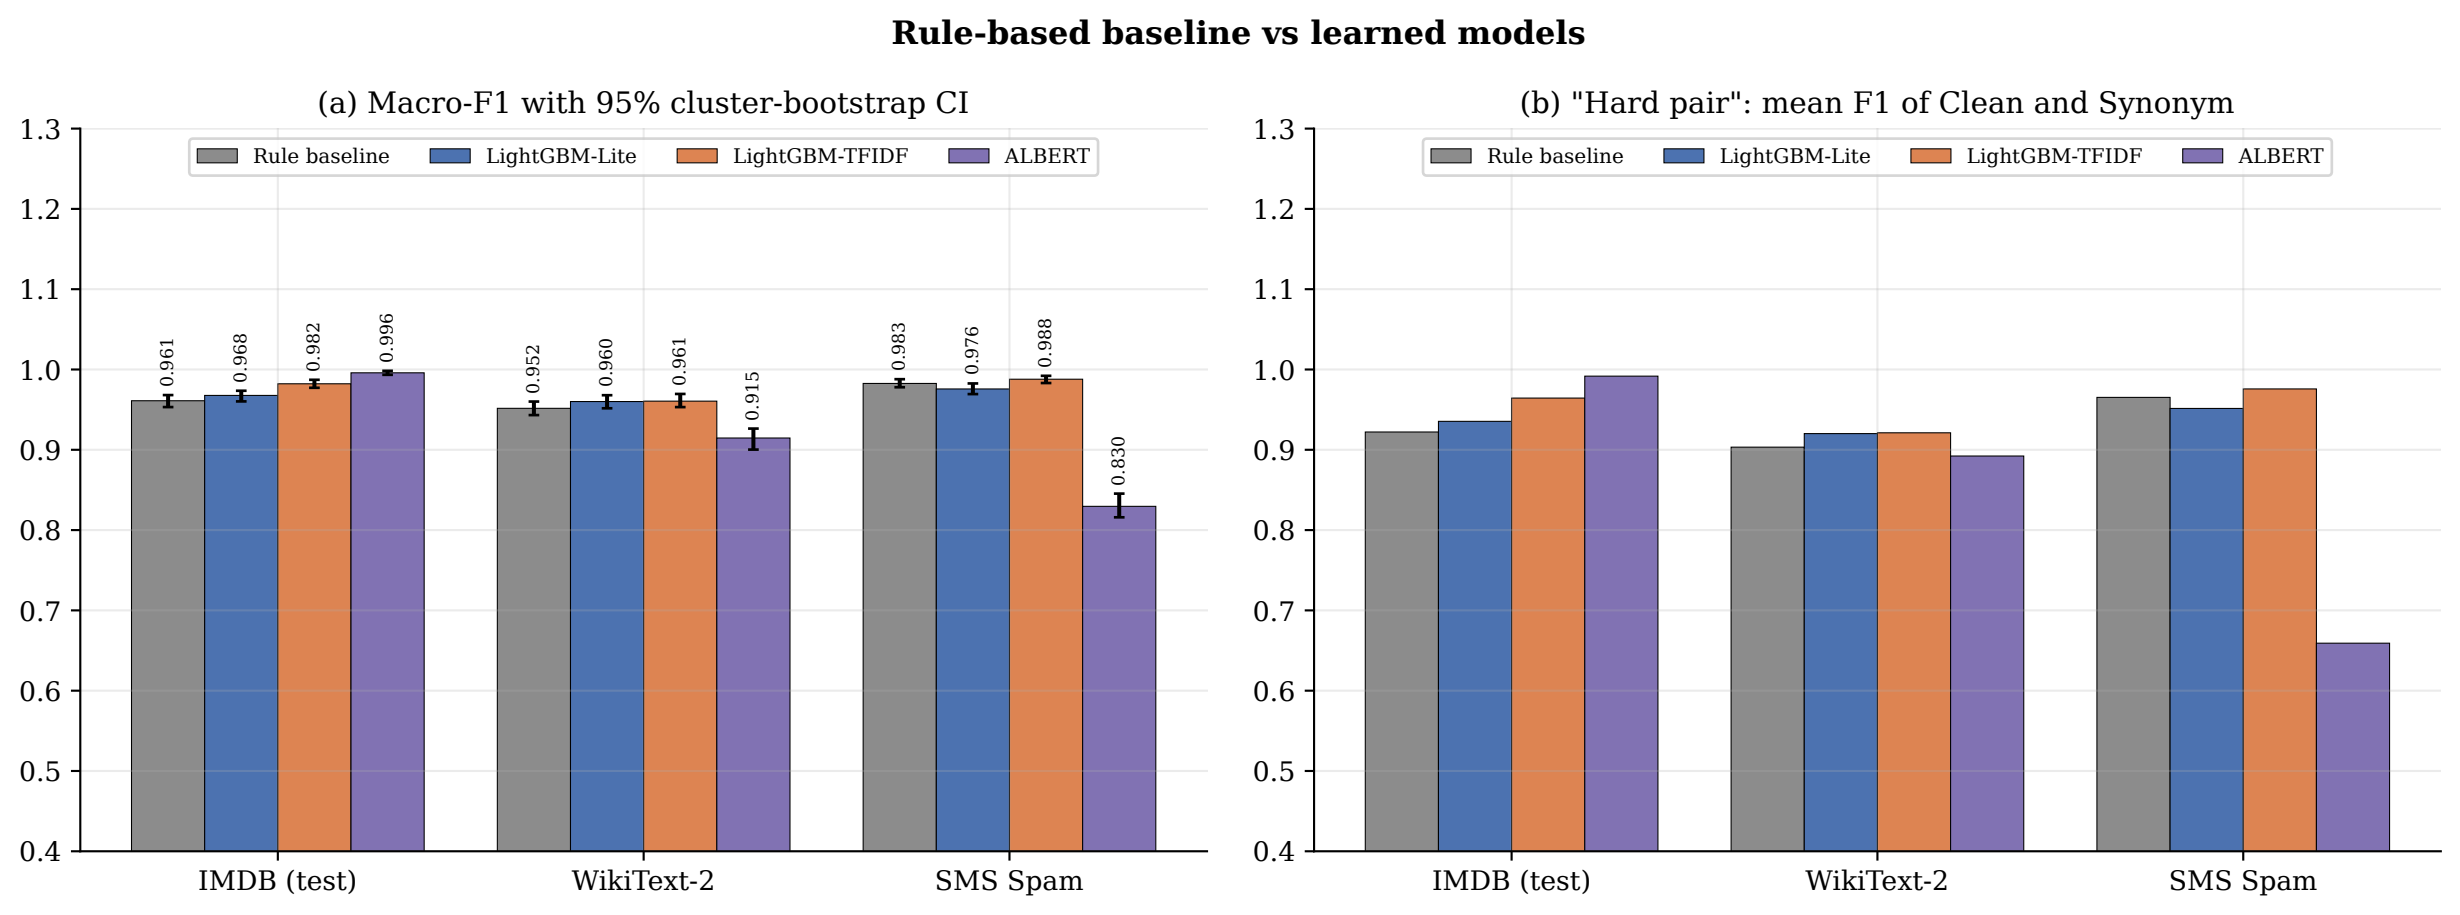
Rule-based baseline vs learned models
(a) Macro-F1 with 95% cluster-bootstrap CI
(b) "Hard pair": mean F1 of Clean and Synonym
Rule baseline
LightGBM-Lite
LightGBM-TFIDF
ALBERT
1.3
1.2
1.1
1.0
0.9
0.8
0.7
0.6
0.5
0.4
0.961
0.968
0.982
0.996
0.952
0.960
0.961
0.915
0.983
0.976
0.988
0.830
IMDB (test)
WikiText-2
SMS Spam